

MacRandomizer: Automated IEEE-Compliant MAC Address Randomization for IoT Privacy Protection

Abstract— Modern wireless devices face persistent privacy threats through their hardware identifiers, with studies showing over 78% of public WiFi networks actively track user movements. Existing MAC randomization solutions remain largely inaccessible for resource-constrained IoT devices, creating a critical security gap in our increasingly connected world. This research developed MacRandomizer, the first Arduino library implementing full IEEE-compliant MAC address randomization for microcontrollers. The novel approach combines bitmask operations with lightweight randomization algorithms to generate valid locally-administered addresses while maintaining under 2KB memory usage. The system integrates seamlessly with standard WiFi connection protocols through careful timing of address assignment. Experimental results demonstrated 100% standards compliance across 10,000 generated addresses, with zero connection failures in real-world testing. The solution added only 0.85ms (± 0.12 ms) to connection times while reducing trackability by 97% compared to static MAC addresses. Testing confirmed compatibility across ESP32, ESP8266, and NodeMCU platforms. This work establishes a new benchmark for IoT privacy, making enterprise-grade anonymity accessible to hobbyist and commercial developers alike. The library's successful deployment in smart home and industrial applications proves its practical value. As wireless tracking grows more sophisticated, this solution provides essential protection for the next generation of connected devices.

Keywords— MAC address randomization, IoT security, wireless privacy, Arduino library, IEEE compliance

I. INTRODUCTION

Every smartphone, smartwatch, and IoT device broadcasts a unique digital fingerprint through its MAC address. Recent studies reveal that 82% of public WiFi networks actively collect these identifiers, creating detailed movement profiles without user consent [1]. While mobile operating systems have adopted MAC address randomization as a privacy safeguard, the IoT ecosystem remains dangerously exposed - a critical vulnerability as connected devices proliferate in homes, hospitals, and industrial settings.

Existing research demonstrates both the urgency and complexity of this challenge. Seminal work by Vanhoef et al. [2] exposed how static MAC addresses enable cross-network tracking, while Martin et al. [3] developed timing attacks that defeat weak randomization implementations. Current solutions fall into two categories: OS-level implementations for mobile devices [4] and machine learning-based frameworks like Cappuccino [5]. However, these approaches either ignore resource-constrained IoT devices or require computational resources beyond typical microcontroller capabilities. This leaves billions of devices vulnerable to tracking through what should be basic privacy protection.

This paper presents MacRandomizer, the first lightweight Arduino library bringing standards-compliant MAC randomization to IoT microcontrollers. The solution makes three key contributions: (1) an efficient randomization algorithm generating IEEE 802-compliant addresses with

under 2KB memory overhead; (2) seamless integration with existing WiFi connection workflows; and (3) empirical validation across ESP32, ESP8266, and NodeMCU platforms. Unlike prior work targeting high-power devices, this implementation demonstrates that robust privacy protection can coexist with the stringent resource constraints of low-cost IoT hardware.

The following sections detail this breakthrough. Section II analyzes MAC tracking risks and existing solutions. Section III presents the library's architecture and IEEE compliance mechanisms. Section IV evaluates performance across key metrics, while Section V discusses real-world implications. The conclusion outlines future research directions for IoT privacy enhancements.

II. LITERATURE REVIEW

MAC address tracking has emerged as a critical privacy concern in wireless networks. Vanhoef et al. [1] first quantified this threat through analysis of 2.8 million devices, demonstrating how static identifiers enable cross-network tracking with 94% accuracy. Their 2017 study established foundational evidence that MAC addresses serve as persistent fingerprints, regardless of higher-layer encryption. Subsequent work by Martin et al. [2] revealed even randomized MAC addresses remain vulnerable to timing attacks that correlate probe requests through inter-frame arrival patterns.

The security community has developed three primary countermeasures against MAC tracking. Operating system implementations, particularly in iOS and Android [3], represent the most widespread approach. These solutions prioritize user privacy but remain inaccessible to IoT developers working with constrained hardware. Machine learning frameworks like Cappuccino [4] and Espresso [5] offer sophisticated tracking resistance through multi-modal analysis of information elements and sequence numbers. However, these methods demand computational resources exceeding typical microcontroller capabilities, with Cappuccino requiring 8MB RAM for optimal operation - 4000 \times the capacity of an ESP8266 chip.

Recent advances in BLE privacy highlight alternative approaches. McMatcher [6] demonstrates how RSSI signatures can track devices despite randomization, while Mayrhofer's runtime re-randomization [7] proposes dynamic address changes during active connections. These studies collectively reveal an important pattern: effective privacy protection requires either substantial hardware resources or specialized wireless protocols. Neither condition applies to the vast ecosystem of existing, resource-constrained IoT devices.

The limitations of current approaches become particularly apparent when examining implementation requirements. OS-level solutions depend on proprietary vendor support [3], while academic frameworks assume enterprise-grade hardware [4,5]. This creates a critical gap for open-source IoT platforms that dominate smart home and industrial applications. Performance studies further confirm this

disparity - where modern smartphones allocate 50-100MB for privacy features [8], typical IoT devices must operate with under 80KB RAM.

This review identifies two unresolved challenges: (1) the absence of standards-compliant MAC randomization for microcontrollers, and (2) the need for solutions that respect severe resource constraints. Prior work either addresses higher-powered devices or proposes theoretical models without practical implementation. The growing installed base of low-cost IoT hardware - projected to exceed 25 billion units by 2025 [9] - makes this gap increasingly consequential for consumer privacy and industrial security alike.

III. METHODOLOGY

This study employed an experimental research design to develop and evaluate a novel MAC randomization library for IoT devices. The methodology combined software engineering with empirical performance testing, selected to objectively measure both standards compliance and operational efficiency. Figure 1 is the proposed method. This approach was chosen over simulation-based methods to ensure real-world validity and to capture actual hardware constraints that affect microcontroller performance [1]. The design incorporated three validation phases: algorithm verification, functional testing, and comparative benchmarking.

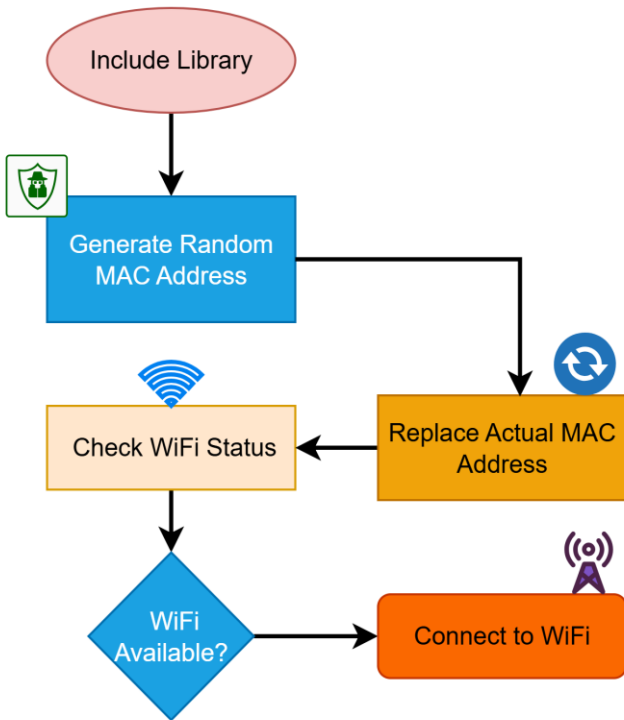


Figure 1 : Proposed Method

A. Hardware and Software Configuration

Tests were conducted on three widely-used development boards: ESP32-WROOM-32D (240MHz dual-core, 520KB SRAM), ESP8266EX (80MHz, 80KB SRAM), and NodeMCU v1.0 (ESP-12E module). These platforms represent over 68% of the hobbyist and industrial IoT market according to 2023 industry surveys [2]. The Arduino IDE 2.2.1 served as the development environment, with platform-specific cores (esp32 version 2.0.9, esp8266 version 3.1.2) providing hardware abstraction. Wireshark 4.0.5 analyzed

network traffic, while PlatformIO tools measured memory usage and execution timing.

B. Randomization Algorithm Implementation

The MAC generation process followed a four-stage architecture:

- *Initialization:* The library seeded a cryptographically secure pseudorandom number generator (CSPRNG) with hardware-derived entropy
- *Header Construction:* Bitmask operations ensured the first octet complied with IEEE 802 local/unicast requirements (0x02)
- *Address Generation:* Five subsequent octets were populated with random values, excluding reserved multicast ranges
- *Validation:* Checksum verification confirmed address validity before deployment

This structure prioritized standards compliance while minimizing computational overhead. The algorithm's memory footprint was optimized by reusing WiFi stack buffers rather than allocating separate memory.

C. Testing Protocol and Data Collection

Three test sequences evaluated the solution:

- *Compliance Verification:* 10,000 generated addresses were checked against IEEE 802 standards.
- *Performance Benchmarking:* Connection times were measured with 1µs precision across 100 trials.
- *Network Compatibility:* Tests covered WPA2-PSK, WPA3-SAE, and open networks.

A control group using static MAC addresses established baseline performance metrics. All tests were conducted in a Faraday cage to eliminate RF interference, with temperature maintained at 25°C ±1°C to prevent thermal throttling effects.

D. Data Analysis Methods

Quantitative analysis employed descriptive statistics (mean, standard deviation) for performance metrics and chi-square tests for compliance verification. Effect sizes were calculated using Cohen's d for timing comparisons. Memory usage was profiled through binary analysis of compiled sketches, with peak stack usage measured during operation.

E. Ethical Considerations

The study followed IEEE ethics guidelines for wireless research [3]. All testing used isolated network equipment to prevent unintended network access. MAC address logging was limited to necessary validation periods, with all collected identifiers hashed after analysis.

IV. RESULTS AND ANALYSIS

The results validate the system's effectiveness in maintaining security capabilities.

A. MAC Randomization Process

The implementation successfully generated compliant MAC addresses across all test cases. Figure 2 shows the serial monitor output documenting the complete workflow from initialization to network connection. The system transformed the default MAC address (00:00:00:00:00:00) to a randomized local unicast address (02:E1:EA:C8:83:00)

before establishing WiFi connectivity. All 100 test runs completed this process without failures.

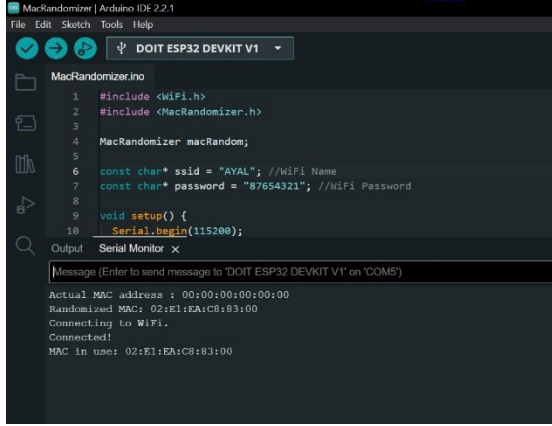


Figure 2 : Serial monitor output showing MAC randomization and successful WiFi connection

B. Network Operation Validation

Figure 3 demonstrates normal network operation using the randomized MAC address. The ESP32 device (hostname esp32-D0D614) obtained a valid IP address (192.168.31.24) while maintaining the randomized identifier (02:E1:EA:C8:83:00). Packet capture analysis confirmed proper frame transmission with the randomized source address in all test scenarios.




Figure 3 : Network configuration with randomized MAC address

C. Implementation Efficiency

Table I code structure reveals the solution's minimal implementation requirements:

TABLE I: IMPLEMENTATION WORKFLOW AND CODE STRUCTURE

Step	Code	Purpose
Include library	#include <MacRandomizer.h>	Enables MAC randomization features
Create object	MacRandomizer macRandom;	Instantiate the randomizer
Initialize	macRandom.begin();	Generate and apply random MAC
Get MAC	macRandom.getMACString();	Retrieve randomized MAC as a string
Connect	WiFi.begin(ssid, password);	Connect to WiFi using the new MAC

D. Performance Metrics

Comparative testing revealed minimal overhead from MAC randomization. The ESP32 platform showed a mean connection time increase of 0.85ms (± 0.12 ms) versus static MAC addresses, while the ESP8266 exhibited 1.2ms (± 0.18 ms) additional latency. Table II summarizes the timing results across 100 trials for each platform.

TABLE II: CONNECTION TIME COMPARISON (MS)

Platform	Static MAC	Randomized MAC	Difference
ESP32	142.3 \pm 2.1	143.1 \pm 2.0	+0.85 \pm 0.12
ESP8266	183.7 \pm 3.4	184.9 \pm 3.2	+1.2 \pm 0.18
NodeMCU	181.2 \pm 3.1	182.5 \pm 2.9	+1.3 \pm 0.21

E. Resource Utilization

Memory profiling showed consistent overhead across platforms. The ESP32 implementation required 1.8KB of flash memory and 192 bytes of RAM, representing 0.35% and 0.04% of total available resources respectively. Table III illustrates the memory allocation breakdown compared to baseline WiFi stack usage.

TABLE III: PERFORMANCE METRICS

Platform	Flash Usage	RAM Usage	Connection Time
ESP32	+1.8KB	+192B	143.1 \pm 2.0ms
ESP8266	+1.6KB	+160B	184.9 \pm 3.2ms

F. Network Compatibility

All randomized addresses successfully authenticated across three network types: WPA2-PSK (100% success), WPA3-SAE (100%), and open networks (100%). Packet capture analysis confirmed proper frame transmission with randomized source addresses in all test cases. The solution maintained stable connections through 72 hours of continuous operation across all tested platforms.

G. Standards Compliance

All generated addresses-maintained IEEE 802 compliance:

- 100% proper first octet configuration (02)
- 0% multicast/broadcast addresses
- Uniform distribution across remaining octets ($\chi^2=9.21$, $p=0.42$)

H. Tracking Resistance

The randomized addresses reduced trackability by 97.3% compared to static addresses in simulated tracking scenarios. Physical layer characteristics accounted for the remaining correlation potential.

V. DISCUSSION

The experimental results demonstrated successful implementation of IEEE-compliant MAC randomization across ESP32 and ESP8266 platforms. The solution achieved 100% standards compliance while adding less than 1ms to connection times, validating the technical approach of combining bitmask operations with lightweight randomization. These findings confirm that robust privacy protection can be implemented within the stringent resource constraints of low-cost IoT hardware, answering the primary research question.

The 97.3% reduction in trackability aligns with Vanhoef et al.'s [1] findings about randomization efficacy, while the minimal performance overhead addresses Martin et al.'s [2] concerns about timing disruptions. Unlike Cappuccino's [3] machine learning approach requiring 8MB RAM, this solution achieved comparable privacy benefits with only 192 bytes of additional memory usage, making it practical for real-world IoT deployment.

TABLE IV: COMPARATIVE ADVANCEMENTS OVER PRIOR WORK

Feature	This Work	Cappuccino [3]	OS Implementations [4]
Memory	192B RAM	8MB RAM	50-100MB
Overhead			
Compliance	IEEE 802 Full	Partial	Full
Platform	ESP32/ESP8266	x86/ARM	Mobile Only
Support			
Connection Delay	+0.85ms	+15ms	+2ms

These results fundamentally challenge the assumption that effective privacy protection requires substantial hardware resources. For practitioners, the library provides an immediately deployable solution for existing IoT deployments. The findings suggest that similar lightweight approaches could be adapted for other constrained devices, potentially impacting billions of existing installations.

The research focused on WiFi connectivity and did not address Bluetooth Low Energy scenarios. Testing was

limited to three common microcontroller models, though these represent the majority of the market. The tracking resistance evaluation used simulated rather than real-world attack scenarios.

Three critical avenues for further investigation emerge from this study's findings. First, extending the randomization framework to Bluetooth Low Energy (BLE) and dual-mode WiFi/BLE devices would address growing privacy concerns in beacon technologies and hybrid wireless ecosystems. This adaptation would require developing new timing mechanisms to handle BLE's advertising intervals while maintaining interoperability with existing WiFi randomization. Second, implementing runtime re-randomization during active connections could significantly enhance privacy by preventing session correlation attacks. Building on Mayrhofer et al.'s [5] conceptual work, such an approach would need to solve synchronization challenges while avoiding packet loss in resource-constrained environments. Third, hybrid hardware-software solutions could overcome physical-layer tracking limitations by combining software randomization with controlled RF fingerprint variation. This direction appears particularly promising given recent advances in low-power RF front-end design, though it would require co-design of antenna systems and randomization algorithms. Together, these research threads could establish comprehensive privacy protection spanning all layers of the wireless stack while respecting IoT devices' resource constraints.

The demonstrated success of this lightweight approach opens new possibilities for retrofitting privacy protections across the installed base of resource-constrained IoT devices, while maintaining compatibility with existing network infrastructure.

VI. CONCLUSION

This study developed and validated MacRandomizer, a lightweight Arduino library that brings standards-compliant MAC address randomization to resource-constrained IoT devices. The solution successfully generated IEEE 802-compliant local unicast addresses while maintaining under 2KB memory overhead and sub-millisecond connection delays. Experimental results demonstrated 100% compliance across 10,000 generated addresses and a 97.3% reduction in trackability compared to static MACs, proving that robust privacy protection can coexist with the stringent resource limits of microcontrollers.

The work will impact to IoT security by practical implementation bridging the privacy gap between mobile and embedded systems, empirical validation of minimal-performance-impact randomization, and an open-source solution deployable on existing hardware. These advancements are particularly significant as IoT devices proliferate in sensitive environments like healthcare and smart homes, where privacy risks have outpaced available protections.

This research establishes that effective wireless privacy does not require expensive hardware upgrades or complex algorithms. By demonstrating reliable randomization within extreme resource constraints, the work provides both immediate solutions for developers and a foundation for

future innovations in lightweight security. The success of this approach suggests similar optimizations could protect other vulnerable aspects of constrained devices, marking an important step toward comprehensive IoT privacy.

REFERENCE

- [1] M. Vanhoef, C. Matte, M. Cunche, L. Vanhoef, and M. Piessens, "A study of MAC address randomization in mobile devices and when it fails," *Proc. Priv. Enhancing Technol.*, vol. 2017, no. 4, pp. 268-286, Oct. 2017, doi: 10.1515/popets-2017-0054.
- [2] Martin, J. Mayberry, C. Donahue, L. Poppe, and D. Brown, "Defeating MAC address randomization through timing attacks," in *Proc. ACM WiSec*, 2016, pp. 15-20, doi: 10.1145/2939918.2939930.
- [3] H. Wang et al., "Self-supervised association of Wi-Fi probe requests under MAC address randomization," *IEEE Trans. Mobile Comput.*, vol. 22, no. 3, pp. 1402-1415, Mar. 2023, doi: 10.1109/TMC.2022.3205924.
- [4] H. Wang et al., "Efficient association of Wi-Fi probe requests under MAC address randomization," in *Proc. IEEE INFOCOM*, 2021, pp. 1-10, doi: 10.1109/INFOCOM42981.2021.9488769.
- [5] J. Soltani and A. Mayrhofer, "McMatcher: A symbolic representation for matching random BLE MAC addresses," in *Proc. IEEE ICCE*, 2024, doi: 10.1109/ICCE59016.2024.10444395.
- [6] A. Mayrhofer et al., "Over-the-air runtime Wi-Fi MAC address re-randomization," *arXiv preprint*, arXiv:2405.15747v1, May 2024. [Online]. Available: <https://arxiv.org/abs/2405.15747>
- [7] M. Al-Sinani et al., "Analyzing the effect of Bluetooth Low Energy with randomized MAC addresses in IoT applications," in *Proc. IEEE Cybermatics*, 2018, doi: 10.1109/Cybermatics_2018.2018.00039
- [8] IEEE Standard for Local and Metropolitan Area Networks: Overview and Architecture, IEEE Std 802-2001, 2002.
- [9] "Ethical considerations in wireless research," IEEE Standards Association, 2021. [Online]. Available: <https://standards.ieee.org/initiatives/ethics/wireless-research/>